



# **TEST PLAN**

## **for**

**ScaleIO Plugin 2.0**  
**Mirantis OpenStack 6.1**

Revision history .....	2
ScaleIO Plugin .....	3
ScaleIO Components .....	3
ScaleIO Cinder and Nova Drivers .....	4
Developer's specification .....	4
Limitations .....	4
Test strategy .....	5
Acceptance criteria .....	5
Test environment, infrastructure and tools .....	5
Product compatibility matrix .....	6
System testing .....	7
Install plugin and deploy environment .....	7
Modifying env with enabled plugin (removing/adding controller nodes) .....	8
Modifying env with enabled plugin (removing/adding compute node) .....	10
Uninstall of plugin with deployed environment .....	11
Uninstall of plugin .....	11
Appendix .....	12

## Revision history

Version	Revision date	Editor	Comment
0.1	28.03.2016	Alexey Morlang (alexey.morlang@emc.com)	Initial version.
0.2			

## ScaleIO Plugin

EMC ScaleIO is a software-only server-based storage area network (SAN) that converges storage and compute resources to form a single-layer, enterprise-grade storage product. ScaleIO storage is elastic and delivers linearly scalable performance. Its scale-out server SAN architecture can grow from a few to thousands of servers. ScaleIO uses servers' direct-attached storage (DAS) and aggregates all disks into a global, shared, block storage. ScaleIO features single-layer compute and storage architecture without requiring additional hardware or cooling/ power/space. Breaking traditional barriers of storage scalability, ScaleIO scales out to hundreds and thousands of nodes and multiple petabytes of storage. The parallel architecture and distributed volume layout delivers a massively parallel system that deliver I/O operations through a distributed system. As a result, performance can scale linearly with the number of application servers and disks, leveraging fast parallel rebuild and rebalance without interruption to I/O. ScaleIO has been carefully designed and implemented with ScaleIO software components so as to consume minimal computing resources. With ScaleIO, any administrator can add, move, or remove servers and capacity on demand during I/O operations. The software responds automatically to any infrastructure change and rebalances data accordingly across the grid nondisruptively. ScaleIO can add capacity on demand, without capacity planning or data migration and grow in small or large increments and pay as you grow, running on any server and with any storage media. ScaleIO natively supports all leading Linux distributions and hypervisors. It works agnostically with any solid-state drive (SSD) or hard disk drive (HDD) regardless of type, model, or speed.

### ScaleIO Components

#### ScaleIO Data Client (SDC)

- Acts as Block Device Driver
- Exposes volumes to applications
- Service must run to provide access to volumes
- Over TCP/IP

#### ScaleIO Data Service (SDS)

- Abstracts storage media
- Contributes to storage pools

- Performs I/O operations

#### ScaleIO Metadata Manager (MDM)

- Not located in the data path
- Provides Monitoring and Configuration management
- Holds cluster-wide component mapping

#### ScaleIO Gateway (GW)

- Provides REST API for Configuration management

### **ScaleIO Cinder and Nova Drivers**

ScaleIO includes a Cinder driver, which interfaces between ScaleIO and OpenStack, and presents volumes to OpenStack as block devices which are available for block storage. It also includes an OpenStack Nova driver, for handling compute and instance volume related operations. The ScaleIO driver executes the volume operations by communicating with the backend ScaleIO MDM through the ScaleIO REST Gateway.

### **Developer's specification**

Is available on [GitHub repository](#).

### **Limitations**

There are following limitations in plugin functionality:

- plugin is currently only compatible with Mirantis 6.1 and 7.0
- plugin supports the only Ubuntu environment
- the only hyper converged environment is supported - there is no separate ScaleIO Storage nodea
- there is no ability to disable completely Cinder-Lvm because of FUEL limitations
- disks for SDS-es should be unallocated, they will be cleaned up
- MDMs and Gateways are deployed together and only onto controller nodes
- there is no ability to separate data network traffic from replication traffic
- there is no fault sets support

## Test strategy

The ScaleIO plugin creates a GUI element to collect the information necessary to deploy and configure EMC ScaleIO in the cluster nodes. The testing strategy is to confirm that all options in the GUI are handled properly and ScaleIO is successfully deployed (if not selected option to use existing ScaleIO) and Cinder and Nova are properly configure to use the ScaleIO cluster as the block storage service and as ephemeral storage service.

### Acceptance criteria

All tests should pass.

### Test environment, infrastructure and tools

The test lab should include 5 nodes. The following designations for the nodes:

- 1) Fuel master node (w/ 50GB Disk, 2 Network interfaces [Mgmt, PXE] )
- 2) OpenStack Controller #1 node
- 3) OpenStack Controller #2 node
- 4) OpenStack Controller #3 node
- 5) OpenStack Compute node

Each node shall have at least 2 CPUs, 4GB RAM, 2x100GB disks, 3 Network interfaces. The 3 interfaces will be used for the following purposes:

- **Admin (PXE) network:** Mirantis OpenStack uses PXE booting to install the operating system, and then loads the OpenStack packages for you.
- **Private, Management and Storage networks:** All of the OpenStack management traffic will flow over this network (“Management” and “Storage” will be separated by VLANs). To re-use the network it will also host the private network. It will be added into each Virtual Machines when they boot. It will therefore be the route where traffic flows in and out of the VM.
- **Private network:** This network is used by OpenStack service nodes and the floating IP address range. Public network must have access to the internet.

### Tool scli how to-s

This is ScaleIO cluster management tool. It is available on all controller nodes after deployment.

1) How to check ScaleIO cluster state with help of 'scli' tool (the tool is available on all controller nodes after deployment):

- read cluster information: `scli --query_cluster`
- login to cluster: `scli --login --mdm_ip <ip_of_master_mdm_from_query_cluster> --username admin --password <password_from_plugin_settings>`
- query information about all SDS: `scli --mdm_ip <ip_of_master_mdm_from_query_cluster> --query_all_sds`
- query information about all SDC: `scli --mdm_ip <ip_of_master_mdm_from_query_cluster> --query_all_sdc`

2) How to remove disconnected SDS and SDC from ScaleIO cluster state with help of 'scli' tool. Interception is presented for SDS, for SDC the actions are the same with just replacing 'sds' words with 'sdc' in cli:

- login to cluster: `scli --login --mdm_ip <ip_of_master_mdm_from_query_cluster> --username admin --password <password_from_plugin_settings>`
- find disconnected SDS/SDC by quering: `scli --mdm_ip <ip_of_master_mdm_from_query_cluster> --query_all_sds`
- remove node from cluster `scli --mdm_ip <ip_of_master_mdm_from_query_cluster> --remove_sds --sds_id <id_of_sds_from_query> --i_am_sure`

## Product compatibility matrix

ScaleIO Plugin version	Compatible Fuel version	OpenStack and OS Version	ScaleIO version
2.0.0	6.1	Juno on Ubuntu14.04	2.0
2.0.0	7.0	Kilo on Ubuntu14.04	2.0

## System testing

### Install plugin and deploy environment

Test Case ID	install_plugin_deploy_env
Steps	<ol style="list-style-type: none"><li>1. Upload plugin to the master fuel node</li><li>2. Install plugin</li><li>3. Ensure that plugin is installed successfully using cli</li><li>4. Create environment with enabled ScaleIO plugin in fuel ui<ol style="list-style-type: none"><li>a. Provide passwords for Admin and for Gateway, e.g. qwe123QWE</li><li>b. Keep protection domain name and storage pool names in default values</li><li>c. Input '/dev/sdb' in Storage devices</li><li>d. Select checkbox 'Controller as SDS'</li></ol></li><li>5. Add 3 node with Controller and Cinder role</li><li>6. Add 1 node with Compute role</li><li>7. Make second disk (sdb) unallocated for all four nodes</li><li>8. Apply network settings<ol style="list-style-type: none"><li>a. IP addresses and assigning networks to interfaces depend on actual network environment of test lab</li></ol></li><li>9. Run network verification</li><li>10. Deploy the cluster</li><li>11. Check ScaleIO cluster state from the primary Controller node</li><li>12. Login to Horizon with the admin user when the OpenStack deployment is finished.<ol style="list-style-type: none"><li>a. Create volume from image</li><li>b. Create flavor with 8GB disks</li><li>c. Run instance with just created flavour</li><li>d. Delete volume and instance created above</li></ol></li><li>13. Run OSTF</li></ol>
Expected Result	<ol style="list-style-type: none"><li>1. Plugin is installed successfully, cluster is created, network verification.</li><li>2. ScaleIO cluster is configured in 'node_3' mode:<ol style="list-style-type: none"><li>1. Two MDMs on the two of controller</li><li>2. One TieBreaker on the one of controller</li><li>3. SDS and SDC on each node and they are in connected state</li><li>4. Volume and instance are created and deleted successfully via Horizon.</li></ol></li><li>3. OSTF are passed.</li></ol>



## Modifying env with enabled plugin (removing/adding controller nodes)

Test Case ID	modify_env_with_plugin_remove_add_controller
Environment	<p>Fuel master node (w/ 50GB Disk, 2 Network interfaces [Mgmt, PXE] )</p> <ol style="list-style-type: none"> <li>1) OpenStack Controller #1 node</li> <li>2) OpenStack Controller #2 node</li> <li>3) OpenStack Controller #3 node</li> <li>4) OpenStack Controller #4 node</li> <li>5) OpenStack Controller #5 node</li> <li>6) OpenStack Compute</li> </ol> <p>Network and disks configuration is the same as described in common section</p>
Steps	<ol style="list-style-type: none"> <li>1. Upload plugin to the master fuel node</li> <li>2. Install plugin</li> <li>3. Ensure that plugin is installed successfully using cli</li> <li>4. Create environment with enabled ScaleIO plugin in fuel ui <ol style="list-style-type: none"> <li>a. Provide passwords for Admin and for Gateway, e.g. qwe123QWE</li> <li>b. Keep protection domain name and storage pool names in default values</li> <li>c. Input '/dev/sdb' in Storage devices</li> <li>d. Select checkbox 'Controller as SDS'</li> </ol> </li> <li>5. Add 3 node with Controller and Cinder role</li> <li>6. Add 1 node with Compute role</li> <li>7. Make second disk (sdb) unallocated for all four nodes</li> <li>8. Apply network settings <ol style="list-style-type: none"> <li>a. IP addresses and assigning networks to interfaces depend on actual network environment of test lab</li> </ol> </li> <li>9. Run network verification</li> <li>10. Deploy the cluster</li> <li>11. Run OSTF</li> <li>12. Remove 1 controller node.</li> <li>13. Check ScaleIO cluster state from the primary Controller node</li> <li>14. Re-deploy cluster</li> <li>15. Check ScaleIO cluster state from the primary Controller node</li> <li>16. Run OSTF</li> <li>17. Add 1 new node with Controller role</li> <li>18. Check ScaleIO cluster state from the primary Controller node</li> <li>19. Re-deploy cluster</li> <li>20. Run OSTF</li> </ol>

Expected Result	<ol style="list-style-type: none"> <li>1. Plugin is installed successfully, cluster is created, network verification.</li> <li>2. ScaleIO cluster is configured in 'node_5' mode: <ol style="list-style-type: none"> <li>1. Three MDMs on the two of controller</li> <li>2. Two TieBreaker on the one of controller</li> <li>3. SDS and SDC on each node and they are in connected state</li> <li>4. Volume and instance are created and deleted successfully via Horizon.</li> </ol> </li> <li>3. OSTF are passed.</li> <li>4. Cluster continues working after removal controller. OSTF passed</li> <li>5. Cluster continues working after adding. OSTF passed</li> </ol>
-----------------	--

## Modifying env with enabled plugin (removing/adding compute node)

Test Case ID	modify_env_with_plugin_remove_add_compute
Steps	<ol style="list-style-type: none"> <li>1. Upload plugin to the master node</li> <li>2. Install plugin</li> <li>3. Ensure that plugin is installed successfully using cli</li> <li>1. Create environment with enabled ScaleIO plugin in fuel ui <ol style="list-style-type: none"> <li>a. Provide passwords for Admin and for Gateway, e.g. qwe123QWE</li> <li>b. Keep protection domain name and storage pool names in default values</li> <li>c. Input '/dev/sdb' in Storage devices</li> <li>d. Select checkbox 'Controller as SDS'</li> </ol> </li> <li>2. Add 3 node with Controller and Cinder role</li> <li>3. Add 1 node with Compute role</li> <li>4. Make second disk (sdb) unallocated for all four nodes</li> <li>5. Apply network settings <ol style="list-style-type: none"> <li>a. IP addresses and assigning networks to interfaces depend on actual network environment of test lab</li> </ol> </li> <li>6. Run network verification</li> <li>7. Deploy the cluster</li> <li>4. Check ScaleIO cluster state from the primary Controller node</li> <li>5. Run OSTF</li> <li>6. Remove 1 compute node</li> <li>7. Re-deploy cluster</li> <li>8. Check ScaleIO cluster state from the primary Controller node, remove disconnected SDS and SDC from ScaleIO cluster via cli</li> <li>9. Run OSTF</li> <li>10. Add 1 compute node</li> <li>11. Re-deploy cluster</li> <li>12. Update controller nodes by running task update_hosts via fuel cli</li> <li>13. Check ScaleIO cluster state from the primary Controller node</li> <li>14. Run OSTF</li> </ol>

Expected Result	<ol style="list-style-type: none"> <li>1. Plugin is installed successfully, cluster is created, network verification.</li> <li>2. ScaleIO cluster is configured in 'node_3' mode: <ol style="list-style-type: none"> <li>1. Two MDMs on the two of controller</li> <li>2. One TieBreaker on the one of controller</li> <li>3. SDS and SDC on each node and they are in connected state</li> <li>4. Volume and instance are created and deleted successfully via Horizon.</li> </ol> </li> <li>3. OSTF are passed.</li> <li>4. Cluster continues working after removal compute node. OSTF passed.</li> <li>5. Cluster continues working after adding compute node. OSTF passed.</li> </ol>
-----------------	---

## Uninstall of plugin with deployed environment

Test Case ID	uninstall_plugin_with_deployed_env
Steps	<ol style="list-style-type: none"> <li>1. install plugin</li> <li>2. deploy environment with enabled plugin functionality</li> <li>3. run ostf</li> <li>4. try to delete plugin and ensure that present in cli alert: "400 Client Error: Bad Request (Can't delete plugin which is enabled for some environment.)"</li> <li>5. remove environment</li> <li>6. remove plugin</li> <li>7. check that it was successfully removed</li> </ol>
Expected Result	<p>Plugin was installed successfully.</p> <p>Alert is present when we trying to delete plugin which is attached to enabled environment.</p> <p>When environment was removed, plugin is removed successfully too.</p>

## Uninstall of plugin

Test Case ID	uninstall_plugin
Steps	<ol style="list-style-type: none"> <li>1. install plugin</li> <li>2. check that it was installed successfully</li> <li>3. remove plugin</li> <li>4. check that it was successfully removed</li> </ol>
Expected Result	Plugin was installed and then removed successfully

## Appendix

№	Resource title
1	<a href="#">ScaleIO Fuel Plugin GitHub repository</a>
2	<a href="#">ScaleIO User Guide</a>
3	<a href="#">ScaleIO OpenStack Information</a>
4	